

# Today

- Integrity Constraints
- ER -> Relational
- Class Participation Exercise

# Integrity Constraints

- ***Integrity Constraint***: a condition specified on the database schema and restricts the data that can be stored.
  - Created when schema is defined
  - Checked when relations are modified
- If a database instance satisfies all the integrity constraints on the schema, then it is a ***legal instance***.
- Changes to the database are not allowed to violate any integrity constraint
- Helps to prevent a whole class of errors and keeps database in a valid state

# Keys

- (From ER) **Key**: minimal set of attributes that uniquely identifies an entity in an entity set
- Similarly, in relational model, a **key constraint** states that a certain minimal subset of fields of a relation is a unique identifier for that tuple
- Any set of fields that uniquely identifies a tuple according to a key constraint is called a **candidate key** for the relation (“key”)
- Out of all candidate keys, can specify a single **primary key** for a relation. Can be used to refer to this tuple elsewhere in the database

# Keys

```
CREATE TABLE Students ( sid    CHAR(20) ,  
                          name  CHAR(30) ,  
                          login  CHAR(20) ,  
                          age    INTEGER,  
                          gpa    REAL,  
                          UNIQUE (name, age),  
                          CONSTRAINT StudentsKey PRIMARY KEY (sid) )
```

- or “PRIMARY KEY (sid)”
- or “sid CHAR(20) PRIMARY KEY”

# NULL Constraints

- By default, most columns (except columns specified as primary keys) can have NULL values
- Can add a constraint that a field can never be NULL
- Could be useful for participation constraints

```
CREATE TABLE Student (  
    sid INT PRIMARY KEY,  
    name VARCHAR(20) NOT NULL)
```

# Foreign Key Constraints

- ***Foreign Key***: Set of fields in one relation that is used to refer to a tuple in another relation.
  - Corresponds to the primary key in the second relation
  - “Logical pointer”

```
CREATE TABLE EnrolledIn (  
    sid INT references Students  
    cid INT references Courses  
    grade CHAR (2)  
    primary key(sid,cid) )
```

# Referential Integrity

- ***Referential Integrity*** requires that foreign keys refer to a valid, existing primary key in another relation
- On INSERT, reject action if foreign key refers to a tuple that doesn't exist
- On DELETE or UPDATE, either
  - Disallow action (default)
  - Cascade changes
  - Set foreign key to default value or NULL

# Referential Integrity

```
CREATE TABLE Enrolled (  
    sid INT  
    cid INT  
    grade CHAR(2)  
    PRIMARY KEY (sid,cid)  
    FOREIGN KEY (sid)  
    REFERENCES Students  
        ON DELETE CASCADE  
        ON UPDATE CASCADE)
```



# Referential Integrity

```
CREATE TABLE Enrolled (  
    sid INT  
    cid INT  
    grade CHAR(2)  
    PRIMARY KEY (sid,cid)  
    FOREIGN KEY (sid)  
    REFERENCES Students  
        ON DELETE CASCADE  
        ON UPDATE SET DEFAULT)
```

# ER -> Relational

- ER model is convenient for representing an initial, high-level database design
- Given ER model, use a standard approach to generate a relational database schema that approximates the ER design
- Approximate because we may not capture all the constraints in the ER design (may be possible in SQL but possibly very expensive)

# Mapping Entity Sets

- Create table for each entity set
- Map attributes to fields
- Declare primary key

# Mapping Relationship Set (no key constraints)

- Create table for relationship set
- Add primary keys of entity sets participating in the relationship as primary keys of the relation
- Map attributes of the relationship to fields

# Mapping Relationship Set (with key constraints)

- Option 1:
  - Create table for relationship set
  - Add primary keys of participating entity sets as fields
  - Map attributes of the relationship to fields
  - Declare primary key using key fields from source entity set (where the arrow is coming from)

# Mapping Relationship Set (with key constraints)

- Option 1:



```
CREATE TABLE manages(  
    employee_id      integer,  
    department_id    integer,  
    PRIMARY KEY (department_id),  
    FOREIGN KEY (employee_id) REFERENCES Employees,  
    FOREIGN KEY (department_id) REFERENCES Departments);
```

# Mapping Relationship Set (with key constraints)

- Option 2
  - Add primary key of target entity as a field in the source

# Mapping Relationship Set (with key constraints)

- Option 2:

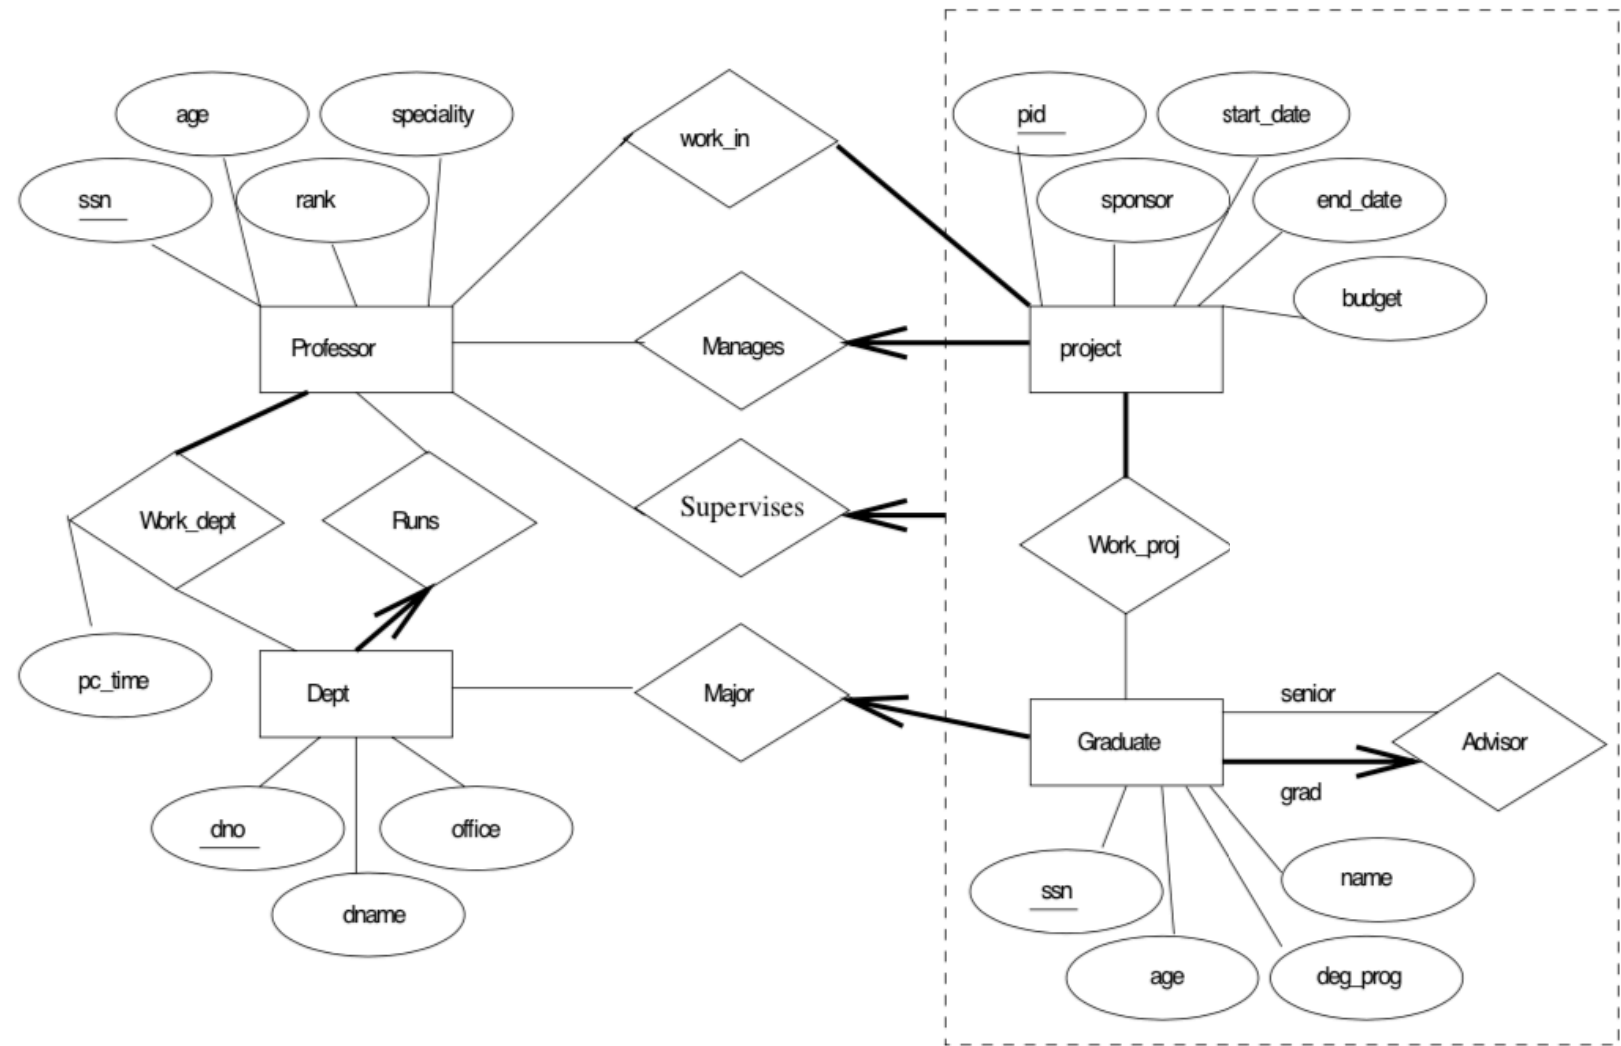


```
CREATE TABLE department(  
    department_id          integer,  
    department_name        varchar(20),  
    employee_id            integer,  
    PRIMARY KEY (department_id),  
    FOREIGN KEY (employee_id) REFERENCES Employees);
```

- Note: if we declare employee\_id as NOT\_NULL, can enforce participation constraint. Cannot easily do this with Option 1.



- Professors have an SSN, a name, an age, a rank, and a research specialty.
- Projects have a project number, a sponsor name (e.g., NSF), a starting date, an ending date, and a budget.
- Graduate students have an SSN, a name, an age, and a degree program
- Each project is managed by one professor
- Each project is worked on by one or more professors
- Professors can manage and/or work on multiple projects.
- Each project is worked on by one or more graduate students
- When graduate students work on a project, a professor must supervise their work on the project.
- Graduate students can work on multiple projects, in which case they will have a (potentially different) supervisor for each one.
- Departments have a department number, a department name, and a main office.
- Departments have a professor (known as the chairman) who runs the department.
- Professors work in one or more departments, and for each department that they work in, a time percentage is associated with their job.
- Graduate students have one major department in which they are working on their degree.
- Each graduate student has another, more senior graduate student (known as a student advisor) who advises him or her on what courses to take.



```
CREATE TABLE Supervises(
  prof_ssn CHAR(10),
  grad_ssn CHAR(10),
  pid      INTEGER,
  PRIMARY KEY (prof_ssn, grad_ssn, pid),
  PRIMARY KEY (grad_ssn, pid),
  FOREIGN KEY (prof_ssn) REFERENCES Professors,
  FOREIGN KEY (grad_ssn) REFERENCES Graduates,
  FOREIGN KEY (pid) REFERENCES Projects);
```

**NOTE:** because of the arrow from the aggregation, each work\_proj is only supervised by one professor. Therefore, The primary key on Supervises should only include grad\_ssn and pid, but NOT prof\_ssn. This is because for each unique work\_proj (defined by a pid and grad\_ssn combination), there should only be one entry in the Supervises table.