

Today

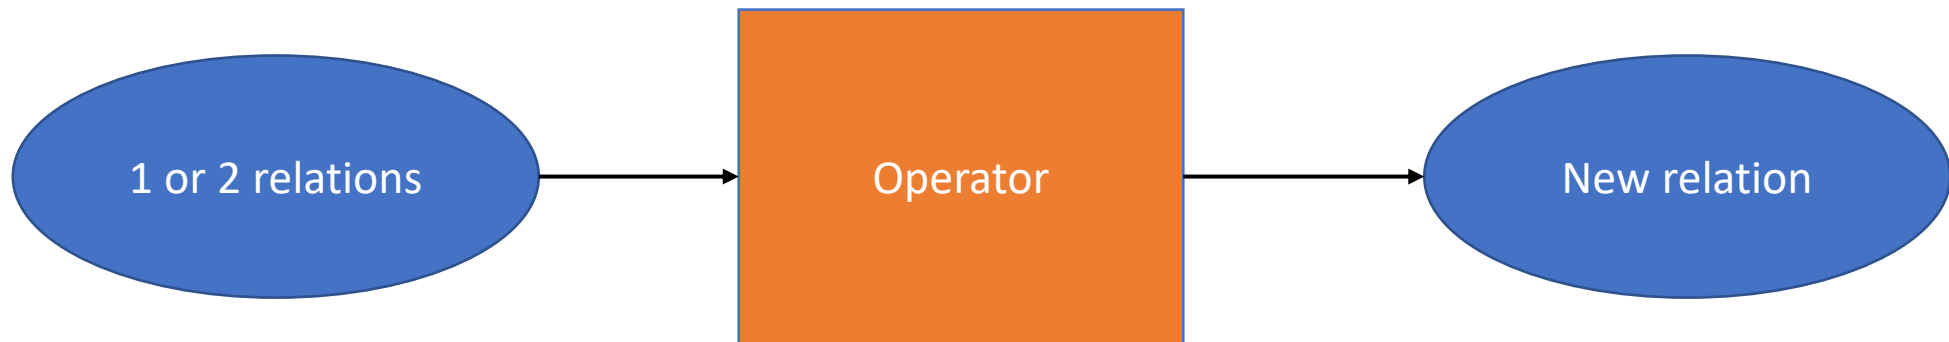
- Review Relational Algebra
- Class Participation Exercise (time permitting)

What is an “Algebra”?

- Mathematical system consisting of:
 - **Operators**: symbols denoting procedures that construct new values from given values (eg. $+$, $-$, $*$, $/$)
 - **Operands**: variables or values from which new values can be constructed. An operand is a value that is acted upon by an operator (eg. $3 + 7$, 3 and 7 are operands)

What is Relational Algebra?

- An algebra whose operands are relations or variables that represent relations
- Input: 1 or 2 instances of relation. Output: an instance of output relation.
- Form basis for other query languages (eg. SQL) and implementation



Operators

- Set Operations:
 - Union
 - Intersection
 - Difference
 - Cross Product (Cartesian Product)
- Selection
- Projection
- Renaming
- Joins
- *Division*

Schema

Sailors(sid: integer, sname: string, rating: integer, age: real)

Boats(bid: integer, bname: string, color: string)

Reserves (sid: integer, bid: integer, day: date)

sid	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.0
58	Rusty	10	35.0

sid	sname	rating	age
28	yuppy	9	35.0
31	Lubber	8	55.0
44	guppy	5	35.0
58	Rusty	10	35.0

Instance S1 of Sailors

Instance S2 of Sailors

sid	bid	day
22	101	10/10/96
58	103	11/12/96

Instance R1 of Reserves

Union

- $R \cup S$ returns an instance of a relation containing all tuples that occur in *either* R or S
- R and S must be ***union-compatible***
 - Same number of fields
 - Corresponding fields, taken in order, must have the same domains
- By convention, inherits names from R
- No duplicates

Union

sid	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.0
58	Rusty	10	35.0

Instance S1 of Sailors

sid	sname	rating	age
28	yuppy	9	35.0
31	Lubber	8	55.0
44	guppy	5	35.0
58	Rusty	10	35.0

Instance S2 of Sailors

sid	sname	rating	age
22	Dustin	7	45.0
28	yuppy	9	35.0
31	Lubber	8	55.0
44	guppy	5	35.0
58	Rusty	10	35.0

S1 U S2

Intersection

- $R \cap S$ returns an instance of a relation containing all tuples that occur in *both* R or S
- R and S must be union-compatible
- By convention, inherits names from R
- No duplicates

Intersection

sid	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.0
58	Rusty	10	35.0

Instance S1 of Sailors

sid	sname	rating	age
28	yuppy	9	35.0
31	Lubber	8	55.0
44	guppy	5	35.0
58	Rusty	10	35.0

Instance S2 of Sailors

sid	sname	rating	age
31	Lubber	8	55.0
58	Rusty	10	35.0

$S1 \cap S2$

Set-Difference

- $R - S$ returns all tuples that occur in R but not S
- Must be union-compatible
- By convention, inherits names from R

Set-Difference

sid	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.0
58	Rusty	10	35.0

Instance S1 of Sailors

sid	sname	rating	age
28	yuppy	9	35.0
31	Lubber	8	55.0
44	guppy	5	35.0
58	Rusty	10	35.0

Instance S2 of Sailors

sid	sname	rating	age
22	Dustin	7	45.0

$S1 - S2$

Cross Product

- $R3 = R1 \times R2$
 - Pair each tuple $t1$ of $R1$ with each tuple $t2$ of $R2$
 - Concatenation $t1, t2$ is a tuple in $R3$
 - Schema of $R3$ is the attributes of $R1$ followed by the attributes of $R2$, in order
 - Beware of attributes having the same name in $R1$ and $R2$: use $R1.A$ and $R2.A$
 - Sometimes called “Cartesian Product”
 - Total number of tuples is equal to $|R1| \times |R2|$ where $|R|$ is the number of tuples in instance relation R

Cross Product

sid	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.0
58	Rusty	10	35.0

Instance S1 of Sailors

sid	bid	day
22	101	10/10/96
58	103	11/12/96

Instance R1 of Reserves

S1.sid	sname	rating	age	R1.sid	bid	day
22	Dustin	7	45.0	22	101	10/10/96
22	Dustin	7	45.0	58	103	11/12/96
31	Lubber	8	55.0	22	101	10/10/96
31	Lubber	8	55.0	58	103	11/12/96
58	Rusty	10	35.0	22	101	10/10/96
58	Rusty	10	35.0	58	103	11/12/96

S1 X R1

Selection σ

- $R1 = \sigma_c(R2)$
 - C is a condition (as in “if” statements) that refers to attributes of R2
 - R1 returns all tuples of R2 that satisfy C
 - Analogous to WHERE clause in SQL

Selection σ

sid	sname	rating	age
28	yuppy	9	35.0
31	Lubber	8	55.0
44	guppy	5	35.0
58	Rusty	10	35.0

Instance S2 of Sailors

sid	sname	rating	age
28	yuppy	9	35.0
58	Rusty	10	35.0

$\sigma_{rating > 8}(S2)$

Projection π

- $R1 = \pi_L(R2)$
 - L is a list of attributes from the schema of $R2$
 - $R1$ is constructed by looking at each tuple of $R2$, extracting the attributes on list L , in the order specified, and creating from those components a tuple from $R1$
 - Eliminate duplicates, if any
 - Analogous to SELECT clause

Projection π

sid	sname	rating	age
28	yuppy	9	35.0
31	Lubber	8	55.0
44	guppy	5	35.0
58	Rusty	10	35.0

Instance S2 of Sailors

age
35.0
55.0

$\pi_{age}(S2)$

sname	rating
yuppy	9
Rusty	10

$\pi_{sname, rating}(\sigma_{rating > 8}(S2))$

Joins

- Conditional Joins: $R3 = R1 \bowtie_c R2$
 - Take the product of $R1 \times R2$
 - Then apply σ_c to the result
 - ie. $R3 = \sigma_c(R1 \times R2)$
 - Result of cross product is usually much larger than result of join, so most implementations do not materialize the underlying cross product

Joins

sid	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.0
58	Rusty	10	35.0

Instance S1 of Sailors

sid	bid	day
22	101	10/10/96
58	103	11/12/96

Instance R1 of Reserves

S1.sid	sname	rating	age	R1.sid	bid	day
22	Dustin	7	45.0	22	101	10/10/96
22	Dustin	7	45.0	58	103	11/12/96
31	Lubber	8	55.0	22	101	10/10/96
31	Lubber	8	55.0	58	103	11/12/96
58	Rusty	10	35.0	22	101	10/10/96
58	Rusty	10	35.0	58	103	11/12/96

$S1 \times R1$

Joins

sid	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.0
58	Rusty	10	35.0

Instance S1 of Sailors

sid	bid	day
22	101	10/10/96
58	103	11/12/96

Instance R1 of Reserves

S1.sid	sname	rating	age	R1.sid	bid	day
22	Dustin	7	45.0	58	103	11/12/96
31	Lubber	8	55.0	58	103	11/12/96

$S1 \bowtie_{S1.sid < R1.sid} R1$

Equijoin

- For join conditions that contain only equalities between attributes of the two relations, there will be a duplicate column for each equality
- In this case an additional projection is done to drop this extra column
- Resulting schema contains fields of R1 followed by fields of R2 that do not appear in the join conditions

Equijoin

sid	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.0
58	Rusty	10	35.0

Instance S1 of Sailors

sid	bid	day
22	101	10/10/96
58	103	11/12/96

Instance R1 of Reserves

S1.sid	sname	rating	age	bid	day
22	Dustin	7	45.0	101	10/10/96
58	Rusty	10	35.0	103	11/12/96

$S1 \bowtie_{S1.sid = R1.sid} R1$

Natural Join

- For a special case of an Equijoin where the join condition contains equalities of all attributes in common between two relations, we can simply remove join condition
- Called a ***Natural Join***
- Denoted $R3 = R1 \bowtie R2$
- Previous $S1 \bowtie_{S1.sid = R1.sid} R1$ is actually a Natural Join since the only common field is sid
- If two relations have no attributes in common, the natural join is simply the cross product

Renaming ρ

- ρ operator allows renaming of relation and attributes
- Useful when fields from different tables have the same name or before certain set operations
- $R' = \rho_{R'(A1 \rightarrow A1', \dots, An \rightarrow An')}(R)$
- R' is a relation with attributes $A1', \dots, An'$ and the same tuples as R

sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	3	25.5
95	Bob	3	63.5

S3 (Sailors)

sid	bid	day
22	101	10/10/98
22	102	10/10/98
22	103	10/8/98
22	104	10/7/98
31	102	11/10/98
31	103	11/6/98
31	104	11/12/98
64	101	9/5/98
64	102	9/8/98
74	103	9/8/98

R2 (Reserves)

bid	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

B1 (Boats)

Q1: Find the names of sailors who have reserved boat 103

sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	3	25.5
95	Bob	3	63.5

S3 (Sailors)

sid	bid	day
22	101	10/10/98
22	102	10/10/98
22	103	10/8/98
22	104	10/7/98
31	102	11/10/98
31	103	11/6/98
31	104	11/12/98
64	101	9/5/98
64	102	9/8/98
74	103	9/8/98

R2 (Reserves)

bid	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

B1 (Boats)

Q1: Find the names of sailors who have reserved boat 103

A1: $\pi_{sname}(\sigma_{bid=103}(Reserves) \bowtie Sailors)$

sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	3	25.5
95	Bob	3	63.5

S3 (Sailors)

sid	bid	day
22	101	10/10/98
22	102	10/10/98
22	103	10/8/98
22	104	10/7/98
31	102	11/10/98
31	103	11/6/98
31	104	11/12/98
64	101	9/5/98
64	102	9/8/98
74	103	9/8/98

R2 (Reserves)

bid	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

B1 (Boats)

Q2: Find the names of sailors who have reserved a red boat.

sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	3	25.5
95	Bob	3	63.5

S3 (Sailors)

sid	bid	day
22	101	10/10/98
22	102	10/10/98
22	103	10/8/98
22	104	10/7/98
31	102	11/10/98
31	103	11/6/98
31	104	11/12/98
64	101	9/5/98
64	102	9/8/98
74	103	9/8/98

R2 (Reserves)

bid	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

B1 (Boats)

Q2: Find the names of sailors who have reserved a red boat.

A2: $\pi_{sname}(\sigma_{color=red}(Boats) \bowtie Reserves \bowtie Sailors)$

sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	3	25.5
95	Bob	3	63.5

S3 (Sailors)

sid	bid	day
22	101	10/10/98
22	102	10/10/98
22	103	10/8/98
22	104	10/7/98
31	102	11/10/98
31	103	11/6/98
31	104	11/12/98
64	101	9/5/98
64	102	9/8/98
74	103	9/8/98

R2 (Reserves)

bid	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

B1 (Boats)

Q3: Find the names of sailors who have reserved a red boat.

sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	3	25.5
95	Bob	3	63.5

S3 (Sailors)

sid	bid	day
22	101	10/10/98
22	102	10/10/98
22	103	10/8/98
22	104	10/7/98
31	102	11/10/98
31	103	11/6/98
31	104	11/12/98
64	101	9/5/98
64	102	9/8/98
74	103	9/8/98

R2 (Reserves)

bid	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

B1 (Boats)

Q3: Find the names of sailors who have reserved a red or a green boat.

A3: $\rho\left(tempboats, (\sigma_{color=red}(Boats)) \cup (\sigma_{color=green}(Boats))\right)$
 $\pi_{sname}(tempboats) \bowtie Reserves \bowtie Sailors$

sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	3	25.5
95	Bob	3	63.5

S3 (Sailors)

sid	bid	day
22	101	10/10/98
22	102	10/10/98
22	103	10/8/98
22	104	10/7/98
31	102	11/10/98
31	103	11/6/98
31	104	11/12/98
64	101	9/5/98
64	102	9/8/98
74	103	9/8/98

R2 (Reserves)

bid	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

B1 (Boats)

Q4: Find the colors of boats reserved by Lubber.

Q5: Find the names of sailors who have reserved a red and a green boat.

sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	3	25.5
95	Bob	3	63.5

S3 (Sailors)

sid	bid	day
22	101	10/10/98
22	102	10/10/98
22	103	10/8/98
22	104	10/7/98
31	102	11/10/98
31	103	11/6/98
31	104	11/12/98
64	101	9/5/98
64	102	9/8/98
74	103	9/8/98

R2 (Reserves)

bid	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

B1 (Boats)

A4: $\pi_{color}(\sigma_{sname=Lubber}(Sailors) \bowtie Reserves \bowtie Boats)$
 A5: $\rho(tempred, \pi_{sid}((\sigma_{color=red}(Boats)) \bowtie Reserves))$
 $\rho(tempgreen, \pi_{sid}((\sigma_{color=green}(Boats)) \bowtie Reserves))$
 $\pi_{sname}(tempred \cap tempgreen) \bowtie Sailors$